

Follow a Swing batch through RLX

01 One policy call

16 worlds: observations (16, 61) → actions (16, 14), values (16, 1)



02 Collect 256 ticks

Observations (256, 16, 61); actions (256, 16, 14)



03 Compute GAE

Rewards, flags, values → advantages and returns (256, 16)



04 Flatten time and world

4096 aligned samples; observations (4096, 61), actions (4096, 14)



05 Shuffle and optimize

4 minibatches of 1024×2 epochs = 8 parameter updates



06 Collect again

Discard old rollout; new actor generates the next 4096 transitions